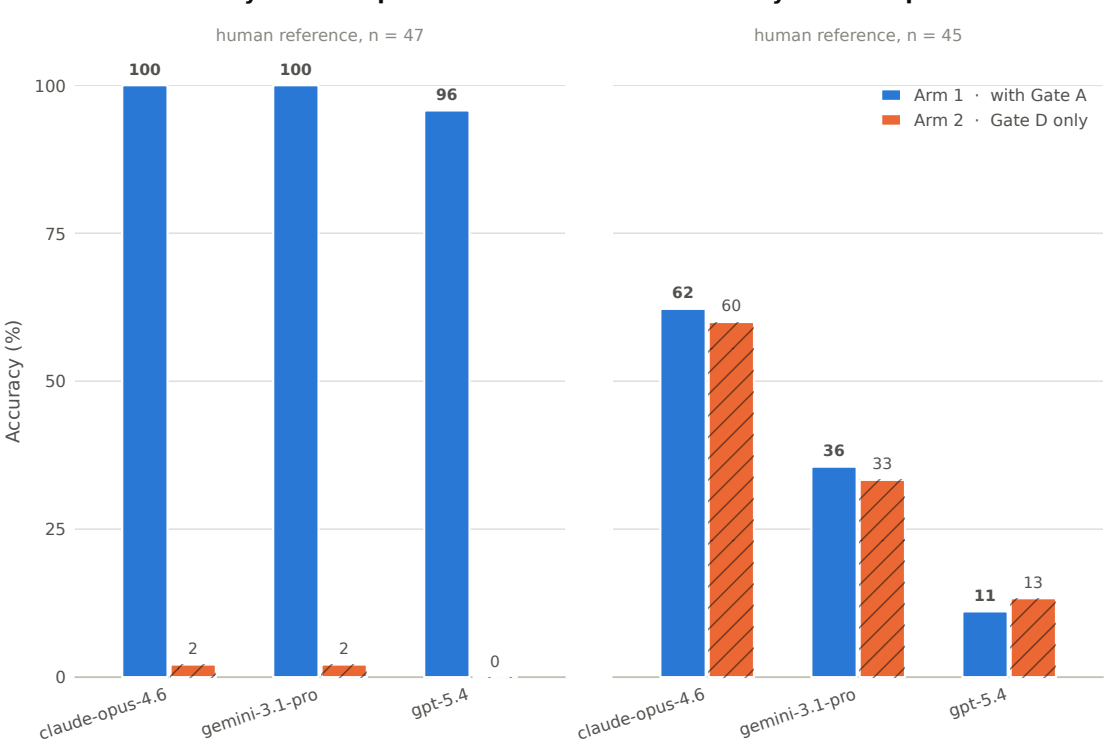

All-97 paired difference (Arm 1 − Arm 2), 95% CI

claude-opus-4.6: +52.6 pp (41.2 to 62.9)
gemini-3.1-pro-preview: +52.6 pp (42.3 to 62.9)
gpt-5.4: +49.5 pp (39.2 to 59.8)

The 5 Gate-A-unassessable cases appear only in the all-97 contrast (47 incorrectly matched + 45 correctly matched + 5 unassessable = 97).

model identifiers accessed through GitHub Copilot on 2026-08-18; not native-provider models